

← 1/40 → *** 5:28:06

Text processing

..



Standing Queries

- **The path from IR to text classification:**
 - You have an information need to monitor, say:
 - Unrest in the Niger delta region
 - You want to rerun an appropriate query periodically to find new news items on this topic
 - You will be sent new documents that are found
 - I.e., it's not ranking but classification (relevant vs. not relevant)
- **Such queries are called standing queries**
 - Long used by “information professionals”
 - A modern mass instantiation is **Google Alerts**
- **Standing queries are (hand-written) text classifiers**

• •

From: Google Alerts

Subject: Google Alert - stanford -neuro-linguistic nlp OR "Natural Language Processing" OR parser OR tagger OR ner OR "named entity" OR segmenter OR classifier OR dependencies OR "core nlp" OR corenlp OR phrasal

Date: May 7, 2012 8:54:53 PM PDT

To: Christopher Manning

Web

3 new results for stanford -neuro-linguistic nlp OR "Natural Language Processing" OR parser OR tagger OR ner OR "named entity" OR segmenter OR classifier OR dependencies OR "core nlp" OR corenlp OR phrasal

[Twitter / Stanford NLP Group: @Robertross If you only n ...](#)

@Robertross If you only need tokenization, java -mx2m edu.stanford.nlp.process.PTBTokenizer file.txt runs in 2MB on a whole file for me.... 9:41 PM Apr 28th ...
twitter.com/stanfordnlp/status/196459102770171905

[\[Java\] LexicalizedParser lp = LexicalizedParser.loadModel\("edu ...](#)

loadModel("edu/stanford/nlp/models/lexparser/englishPCFG.ser.gz"); String[] sent = { "This", "is", "an", "easy", "sentence", "." }; Tree parse = lp.apply(Arrays.
pastebin.com/az14R9nd

[More Problems with Statistical NLP || kuro5hin.org](#)

Tags: nlp, ai, coursera, stanford, nlp-class, cky, nltk, reinventing the wheel, ... Programming Assignment 6 for Stanford's nlp-class is to implement a CKY parser .
www.kuro5hin.org/story/2012/5/5/11011/68221

Tip: Use quotes ("like this") around a set of words in your query to match them exactly. [Learn more](#).

[Delete](#) this alert.

[Create](#) another alert.

[Manage](#) your alerts.

••



Spam filtering Another text classification task

From: "" <takworlld@hotmail.com>

Subject: real estate is the only way... gem oalvgkay

Anyone can buy real estate with no money down

Stop paying rent TODAY !

There is no need to spend hundreds or even thousands for similar courses

I am 22 years old and I have already purchased 6 properties using the methods outlined in this truly INCREDIBLE ebook.

Change your life NOW !

Click Below to order:

<http://www.wholesaledaily.com/sales/nmd.htm>

..



USC **Viterbi**
School of Engineering


Categorization/Classification

- Given:
 - A representation of a document d
 - Issue: how to represent text documents.
 - Usually some type of high-dimensional space – bag of words
 - A fixed set of classes:
$$C = \{c_1, c_2, \dots, c_J\}$$
- Determine:
 - The category of d by generating a classification function, say $\gamma(d)$
 - We want to build classification functions (“classifiers”).

Copyright Ellis Horowitz, 2011-2015

5

••



..




USC **Viterbi**
School of Engineering

Classification Methods (1)

- **Manual classification**

- Used by the original Yahoo! Directory
- Looksmart, about.com, ODP, PubMed
- Accurate when job is done by experts
- Consistent when the problem size and team is small
- Difficult and expensive to scale
 - Means we need automatic classification methods for big problems



Copyright Ellis Horowitz, 2011-2015

7


..



USC **Viterbi**
School of Engineering

Classification Methods (2)

- **Hand-coded rule-based classifiers**
 - One technique used by news agencies, intelligence agencies, etc.
 - Widely deployed in government and enterprises
 - Vendors provide “IDE” for writing such rules



Copyright Ellis Horowitz, 2011-2015


8

..



Classification Methods (2)

- **Hand-coded rule-based classifiers**
 - Commercial systems have complex query languages
 - Accuracy is can be high if a rule has been carefully refined over time by a subject expert
 - Building and maintaining these rules is expensive

A faint, semi-transparent watermark of a person holding a book is visible in the background of the slide.

Copyright Ellis Horowitz, 2011-2015


9

..



Classification Methods (3): Supervised learning

- Given:
 - A document d
 - A fixed set of classes:
$$C = \{c_1, c_2, \dots, c_J\}$$
 - A training set D of documents each with a label in C
- Determine:
 - A learning method or algorithm which will enable us to learn a classifier γ
 - For a test document d , we assign it the class
$$\gamma(d) \in C$$




..



Classification Methods (3)

- **Supervised learning**
 - **Naive Bayes (simple, common)**
 - **k-Nearest Neighbors (simple, powerful)**
 - **Support-vector machines (newer, generally more powerful)**
 - ... plus many other methods
 - **No free lunch: requires hand-classified training data**
 - **But data can be built up (and refined) by amateurs**
- **Many commercial systems use a mixture of methods**



Copyright Ellis Horowitz, 2011-2015

11

..



USC **Viterbi**
School of Engineering

The bag of words representation

Y(

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet.

)=C

Copyright Ellis Horowitz, 2011-2015

12

..



The bag of words representation

Y(

great	2
love	2
recommend	1
laugh	1
happy	1
...	...

)=C

Copyright Ellis Horowitz, 2011-2015


13

..



Features

- **Supervised learning classifiers can use any sort of feature**
 - URL, email address, punctuation, capitalization, dictionaries, network features
- **In the simplest bag of words view of documents**
 - We use only word features
 - we use all of the words in the text (not a subset)



Copyright Ellis Horowitz, 2011-2015

14


..



USC **Viterbi**
School of Engineering

Feature Selection: Why?

- **Text collections have a large number of features**
 - 10,000 – 1,000,000 unique words ... and more
- **Selection may make a particular classifier feasible**
 - Some classifiers can't deal with 1,000,000 features
- **Reduces training time**
 - Training time for some methods is quadratic or worse in the number of features
- **Makes runtime models smaller and faster**
- **Can improve generalization (performance)**
 - Eliminates noise features
 - Avoids overfitting



Copyright Ellis Horowitz, 2011-2015

15


..



USC **Viterbi**
School of Engineering

Feature Selection: Frequency

- **The simplest feature selection method:**
 - Just use the most common terms
 - No particular foundation
 - But it make sense why this works
 - They are the words that can be well-estimated and are most often available as evidence
 - In practice, this is often 90% as good as better methods
 - Smarter feature selection – future lecture



Copyright Ellis Horowitz, 2011-2015


16

..



SpamAssassin

- **Naïve Bayes has found a home in spam filtering**
 - **Paul Graham's A Plan for Spam**
 - <http://www.paulgraham.com/spam.html>
 - **Widely used in spam filters**
 - **But many features beyond words:**
 - **black hole lists, etc.**
 - **particular hand-crafted text patterns**



Copyright Ellis Horowitz, 2011-2015

17

Here is Paul's page, on his spam filtering expts.

Here is SpamAssassin, and this is an old page with results of tests performed.

..



USC **Viterbi**
School of Engineering

Naive Bayes is Not So Naive

- Very fast learning and testing (basically just count words)
- Low storage requirements
- Very good in domains with many equally important features
- More robust to irrelevant features than many learning methods

Irrelevant features cancel each other without affecting results

Copyright Ellis Horowitz, 2011-2015


19

..



Evaluating Categorization

- **Measures: precision, recall, F1, classification accuracy**
- **Classification accuracy:** r/n where n is the total number of test docs and r is the number of test docs correctly classified



Copyright Ellis Horowitz, 2011-2015

22

..



USC **Viterbi**
School of Engineering

WebKB Experiment (1998)

- Classify webpages from CS departments into:
 - student, faculty, course, project
- Train on ~5,000 hand-labeled web pages
 - Cornell, Washington, U.Texas, Wisconsin
- Crawl and classify a new site (CMU) using Naïve Bayes
- Results

	Student	Faculty	Person	Project	Course	Departmt
Extracted	180	66	246	99	28	1
Correct	130	28	194	72	25	1
Accuracy:	72%	42%	79%	73%	89%	100% 23

Copyright Ellis Horowitz, 2011-2015


..



USC **Viterbi**
School of Engineering

Recall: Vector Space Representation

- **Each document is a vector, one component for each term (= word).**
- **Normally normalize vectors to unit length.**
- **High-dimensional vector space:**
 - Terms are axes
 - 10,000+ dimensions, or even 100,000+
 - Docs are vectors in this space
- **How can we do classification in this space?**




..



Classification Using Vector Spaces

- In vector space classification, training set corresponds to a labeled set of points (equivalently, vectors)
- Premise 1: Documents in the same class form a contiguous region of space
- Premise 2: Documents from different classes don't overlap (much)
- Learning a classifier: build surfaces to delineate classes in the space

A faint, semi-transparent watermark or background image of a person sitting and reading a book, visible behind the text area.


Copyright Ellis Horowitz, 2011-2015

26

••




Documents in a Vector Space




Copyright Ellis Horowitz 2011-2012

27

..



••



..



Definition of centroid

$$\vec{\mu}(c) = \frac{1}{|D_c|} \sum_{d \in D_c} \vec{v}(d)$$

- Where D_c is the set of all documents that belong to class c and $v(d)$ is the vector space representation of d .
- Note that centroid will in general not be a unit vector even when the inputs are unit vectors.

..



Rocchio classification

- Rocchio forms a simple representative for each class: the centroid/prototype
- Classification: nearest prototype/centroid
- It does not guarantee that classifications are consistent with the given training data

Why not?

Copyright Ellis Horowitz, 2011-2015

31 31


In Rocchio classification, centroids (one for each term group) are used to specify regions; lines/planes/hyperplanes between centroids produce convex **Voronoi** regions. The new/incoming term's closest centroid is used to classify the term.

..



Rocchio classification

- A simple form of Fisher's linear discriminant
- Little used outside text classification
 - It has been used quite effectively for text classification
 - But in general worse than Naïve Bayes
- Again, cheap to train and test documents



..




USC **Viterbi**
School of Engineering

***k* Nearest Neighbor Classification**

- ***kNN = k* Nearest Neighbor**


- **To classify a document d :**
- **Define k -neighborhood as the k nearest neighbors of d**
- **Pick the majority class label in the k -neighborhood**
- **For larger k can roughly estimate $P(c|d)$ as $\#(c)/k$**



Copyright Ellis Horowitz, 2011-2015

35 35

..




..



USC **Viterbi**
School of Engineering

Nearest-Neighbor Learning

- Learning: just store the labeled training examples D
- Testing instance x (*under 1NN*):
 - Compute similarity between x and all examples in D .
 - Assign x the category of the most similar example in D .
- Does not compute anything beyond storing the examples
- Also called:
 - Case-based learning
 - Memory-based learning
 - Lazy learning
- Rationale of kNN: contiguity hypothesis



Copyright Ellis Horowitz, 2011-2015

37


37

..



k Nearest Neighbor

- Using only the closest example (1NN) subject to errors due to:
 - A single atypical example.
 - Noise (i.e., an error) in the category label of a single training example.
- More robust: find the k examples and return the majority category of these k
- k is typically odd to avoid ties; 3 and 5 are most common



..



Nearest Neighbor with Inverted Index

- Naively finding nearest neighbors requires a linear search through $|D|$ documents in collection
- But determining k nearest neighbors is the same as determining the k best retrievals using the test document as a query to a database of training documents.
- Use standard vector space inverted index methods to find the k nearest neighbors.
- **Testing Time:** $O(B|V_t|)$ where B is the average number of training documents in which a test-document word appears.
 - Typically $B \ll |D|$

Copyright Ellis Horowitz, 2011-2015


39 39

..



kNN: Discussion

- No feature selection necessary
- No training necessary
- Scales well with large number of classes
 - Don't need to train n classifiers for n classes
- Classes can influence each other
 - Small changes to one class can have ripple effect
- Done naively, very expensive at test time
- In most cases it's more accurate than NB or Rocchio




..



USC **Viterbi**
School of Engineering

Rocchio Anomaly

- Prototype models have problems with polymorphic (disjunctive) categories.



Copyright Ellis Horowitz, 2011-2015


42 42

..



3 Nearest Neighbor vs. Rocchio

- **Nearest Neighbor tends to handle polymorphic categories better than Rocchio/NB.**



Copyright Ellis Horowitz, 2011-2015


43

..



Bias vs. capacity – notions and terminology

- Consider asking a botanist: Is an object a tree?
 - Too much *capacity*, low *bias*
 - Botanist who memorizes
 - Will always say “no” to new object (e.g., different # of leaves)
 - Not enough capacity, high bias
 - Lazy botanist
 - Says “yes” if the object is green
 - You want the middle ground




Copyright Ellis Horowitz, 2011-2015 (Example due to C. Burges) 44

..



kNN vs. Naive Bayes

- Bias/Variance tradeoff
 - Variance \approx Capacity
- kNN has high variance and low bias.
 - Infinite memory
- Rocchio/NB has low variance and high bias.
 - Linear decision surface between classes



Copyright Ellis Horowitz, 2011-2015


45 45

..



Summary: Representation of Text Categorization Attributes

- Representations of text are usually very high dimensional
 - “The curse of dimensionality”
- High-bias algorithms should generally work best in high-dimensional space
 - They prevent overfitting
 - They generalize more
- For most text categorization tasks, there are many relevant features and many irrelevant ones



Copyright Ellis Horowitz, 2011-2015

47

47

..



Which classifier do I use for a given text classification problem?

- Is there a learning method that is optimal for all text classification problems?
- No, because there is a tradeoff between bias and variance.
- Factors to take into account:
 - How much training data is available?
 - How simple/complex is the problem? (linear vs. nonlinear decision boundary)
 - How noisy is the data?
 - How stable is the problem over time?
 - For an unstable problem, it's better to use a simple and robust classifier.

